

# Audio\_classification

April 16, 2023

## 1 Music genre Classification

Dataset: The GTZAN dataset is the most-used public dataset for evaluation in machine listening research for music genre recognition (MGR). The files were collected in 2000-2001 from a variety of sources including personal CDs, radio, microphone recordings, in order to represent a variety of recording conditions (<http://marsyas.info/downloads/datasets.html>). The dataset include: - genres original - A collection of 10 genres with 100 audio files each, all having a length of 30 seconds (the famous GTZAN dataset, the MNIST of sounds) - images original - A visual representation for each audio file. One way to classify data is through neural networks. Because NNs (like CNN, what we will be using today) usually take in some sort of image representation, the audio files were converted to Mel Spectrograms to make this possible. - 2 CSV files - Containing features of the audio files. One file has for each song (30 seconds long) a mean and variance computed over multiple features that can be extracted from an audio file. The other file has the same structure, but the songs were split before into 3 seconds audio files (this way increasing 10 times the amount of data we fuel into our classification models). With data, more is always better.

Our project will use GTZAN dataset for music genre classification. The process includes the following 5 parts. 1. Visualize audio dataset: soundwave visualization, Fourier transform, Spectrogram, Mel Spectrogram, Harmonics and percussive, MFCC 2. EDA: Correlation among features & PCA 4. Classification through feature data: Gtzan dataset have extracted import features as a new dataset, such as chroma\_stft\_mean, chroma\_stft\_var, rms\_mean, rms\_var, spectral\_centroid\_mean. We can use it as a feature dataset for classification. 5. Classification through audio images: Beside classifying through features, we also process the image data from audio files and perform classification. In the end, we compare the performance between classifying through feature data and audio images.

### 1.1 1. Visualize Audio Data

```
[ ]: # Usual Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import sklearn

# Librosa (the mother of audio files)
import librosa
```

```
import librosa.display
import IPython.display as ipd
```

```
[ ]: import os
path = 'GTZAN/'
print(list(os.listdir(f'{path}/genres_original/')))
```

```
['pop', 'metal', 'disco', 'blues', 'reggae', 'classical', 'rock', 'hiphop',
'country', 'jazz']
```

```
[ ]: # Importing 1 file
y_pop, sr_pop = librosa.load(f'{path}/genres_original/pop/pop.00036.wav')

print('y:', y_pop, '\n')
print('y shape:', np.shape(y_pop), '\n')
print('Sample Rate (KHz):', sr_pop, '\n')

# Verify length of the audio
print('Check Len of Audio:', 661504/22050)
```

```
y: [-0.19229126 -0.08969116  0.02322388 ... -0.16653442 -0.0585022
 0.01345825]
```

```
y shape: (661504,)
```

```
Sample Rate (KHz): 22050
```

```
Check Len of Audio: 30.00018140589569
```


The sample rate of 22050 Hz was commonly used in early digital audio formats, such as the Audio Compact Disc (CD), which was introduced in the 1980s. CD audio is sampled at a rate of 44.1 kHz, which is twice the sample rate of 22050 Hz.

The reason why 44.1 kHz was chosen as the standard for CD audio is because it is capable of accurately representing audio frequencies up to 20 kHz, which is the upper limit of human hearing.

### 1.1.1 2D Representation: Sound Waves

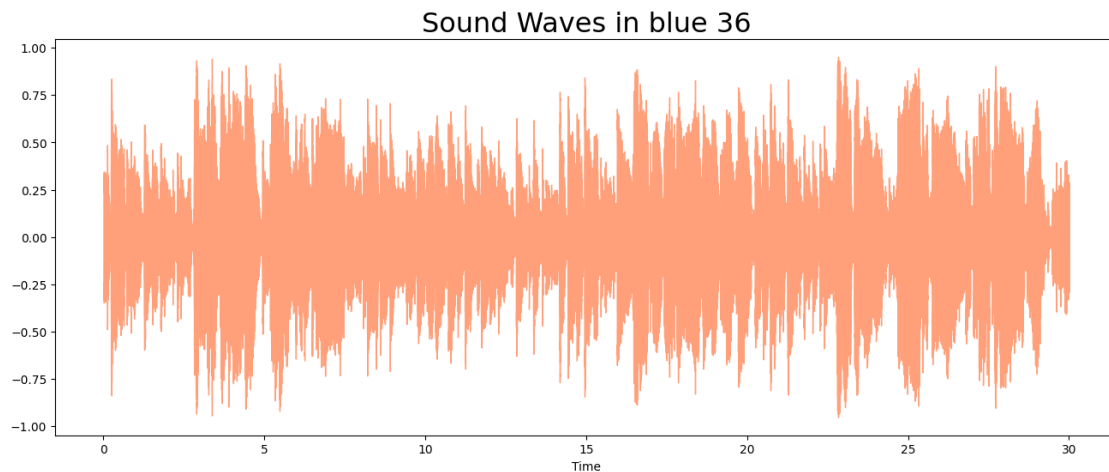
Note on waveshow: x-axis represents time and the y-axis represents the amplitude of the audio signal.

```
[ ]: plt.figure(figsize = (16, 6))
librosa.display.waveshow(y=y_pop, sr=sr_pop, color='#FFA07A')
plt.title("Sound Waves in Pop 36", fontsize = 23);
```




```
[ ]: y_blue, sr_blue = librosa.load(f'{path}/genres_original/blues/blues.00036.wav')

plt.figure(figsize = (16, 6))
librosa.display.waveshow(y=y_blue, sr=sr_blue, color='#FFA07A')
plt.title("Sound Waves in blue 36", fontsize = 23);
```




```
[ ]: y_classical, sr_classical = librosa.load(f'{path}/genres_original/classical/
↪classical.00036.wav')

plt.figure(figsize = (16, 6))
librosa.display.waveshow(y=y_classical, sr=sr_classical, color='#FFA07A')
plt.title("Sound Waves in Classical 36", fontsize = 23);
```



```
[ ]: y_country, sr_country = librosa.load(f'{path}/genres_original/country/country.
    ↪00036.wav')

plt.figure(figsize = (16, 6))
librosa.display.waveshow(y=y_country, sr=sr_country, color='#FFA07A')
plt.title("Sound Waves in Country 36", fontsize = 23);
```



### 1.1.2 Fourier Transform

Function that gets a signal in the time domain as input, and outputs its decomposition into frequencies. Transform both the y-axis (frequency) to log scale, and the “color” axis (amplitude) to Decibels, which is approx. the log scale of amplitudes.

```
[ ]: # Default FFT window size
n_fft = 2048 # FFT window size
hop_length = 512 # number audio of frames between STFT columns (looks like a
↳ good default)


# Short-time Fourier transform (STFT)
D = np.abs(librosa.stft(y_pop, n_fft = n_fft, hop_length = hop_length))

print('Shape of D object:', np.shape(D))
```

Shape of D object: (1025, 1293)

The purpose of STFT is to add the time dimension back by breaking down signals into windows. (Window length: the time length of the window, hope length: the extent of window overlapping. The longer the hop length, the higher the freq resolutions (larger segments of the audio signals), but the lower the time resolution (few frames per unit of time) We then perform a Fourier transformation of the window to understand the pattern of the sounds.

```
[ ]: plt.figure(figsize = (16, 6))
plt.plot(D);
```




The resulting plot will show the frequency content of the audio signal over time. The x-axis of the plot represents time, with each point on the axis corresponding to a different time frame. The y-axis represents the frequency content of the signal, with each point on the axis corresponding to a different frequency band.

### 1.1.3 Spectrogram

The horizontal axis of a spectrogram represents time, while the vertical axis represents frequency. The intensity of the spectrogram is represented by color or grayscale, with brighter colors indicating higher intensity or amplitude at a given frequency.

```
[ ]: # Convert an amplitude spectrogram to Decibels-scaled spectrogram.
DB = librosa.amplitude_to_db(D, ref = np.max)


# Creating the Spectrogram
plt.figure(figsize = (16, 6))
librosa.display.specshow(DB, sr = sr_pop, hop_length = hop_length, x_axis = 'time',
    y_axis = 'log',
    cmap = 'viridis')
plt.colorbar();
```



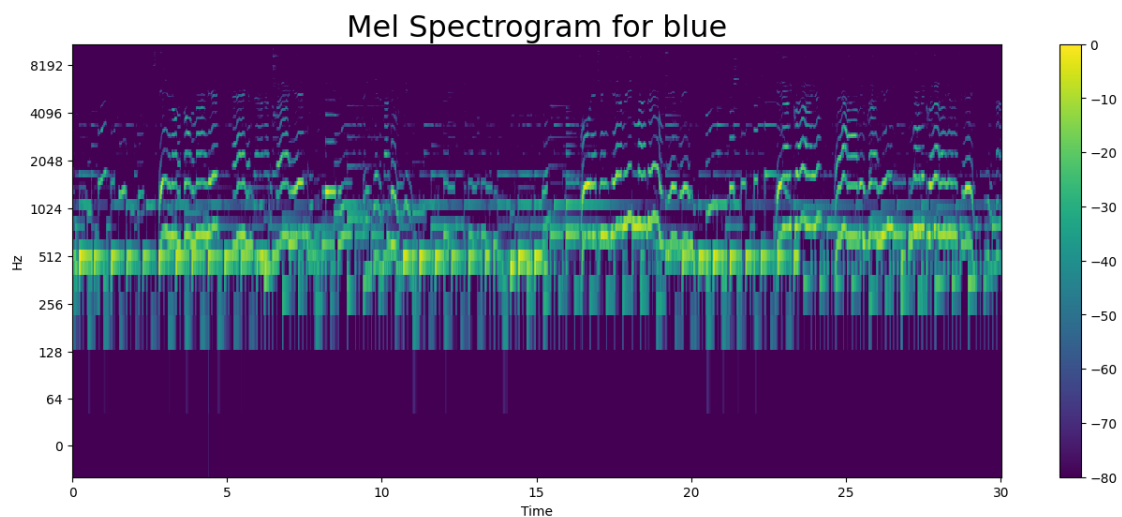
### 1.1.4 Mel Spectrogram

Mel Scale - frequency for humans (Reason: human perceive the sound distance unlinearly)

```
[ ]: S = librosa.feature.melspectrogram(y = y_pop, sr=sr_pop)
S_DB = librosa.amplitude_to_db(S, ref=np.max)
plt.figure(figsize = (16, 6))
librosa.display.specshow(S_DB, sr=sr_pop, hop_length=hop_length, x_axis = 'time',
    y_axis = 'log',
    cmap = 'viridis');
plt.colorbar();
plt.title("Mel Spectrogram for pop", fontsize = 23);
```




```
[ ]: S = librosa.feature.melspectrogram(y = y_blue, sr=sr_blue)
S_DB = librosa.amplitude_to_db(S, ref=np.max)
plt.figure(figsize = (16, 6))
librosa.display.specshow(S_DB, sr=sr_blue, hop_length=hop_length, x_axis = 'time', y_axis = 'log',
                          cmap = 'viridis');
plt.colorbar();
plt.title("Mel Spectrogram for blue", fontsize = 23);
```



```
[ ]: S = librosa.feature.melspectrogram(y = y_classical, sr=sr_classical)
S_DB = librosa.amplitude_to_db(S, ref=np.max)
```

```
plt.figure(figsize = (16, 6))
librosa.display.specshow(S_DB, sr=sr_classical, hop_length=hop_length, x_axis = 'time', y_axis = 'log',
                        cmap = 'viridis');
plt.colorbar();
plt.title("Mel Spectrogram for classical", fontsize = 23);
```



Observations: Pop music have higher amplitudes. Blue music have smaller range of frequency. Classical music seems to have the most dynamic frequency.

### 1.1.5 Audio Features

Zero Crossing Rate: Zero Crossing Rate (ZCR) is a measure of the number of times a signal crosses the zero-axis (amplitude) in a given time frame.

```
[ ]: # Total zero_crossings for pop
zero_crossings = librosa.zero_crossings(y_pop, pad=False)
print(sum(zero_crossings))

# Total zero_crossings for blue
zero_crossings = librosa.zero_crossings(y_blue, pad=False)
print(sum(zero_crossings))

# Total zero_crossings in our 1 song
zero_crossings = librosa.zero_crossings(y_classical, pad=False)
print(sum(zero_crossings))
```

76229

24463

39405



### 1.1.6 Harmonics and percussive

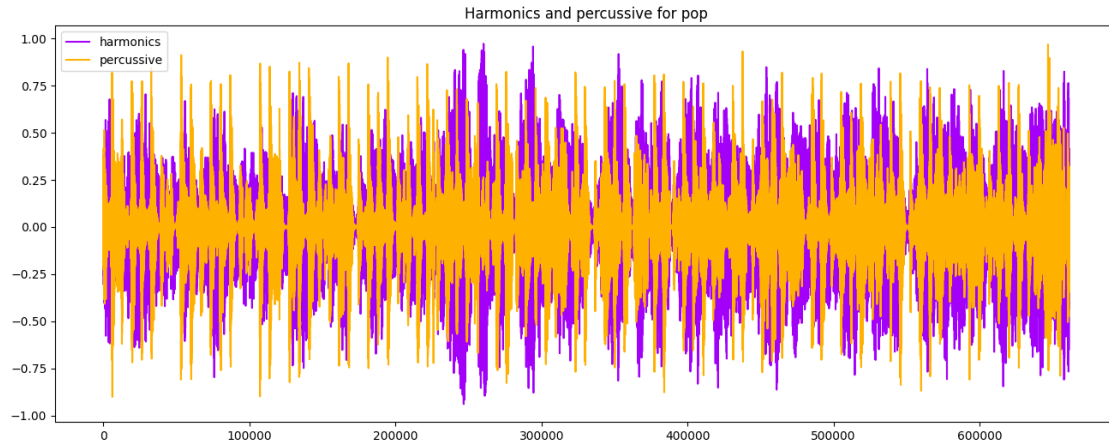
- Harmonics: These are pitch sounds that enable us to hear melodies.
- Percussive sound: This is more like something originating from an instrument onset, like a beat on a drum.

The x-axis represents time and the y-axis represents the amplitude of the signal.

```
[ ]: y_harm, y_perc = librosa.effects.hpss(y_pop)

plt.figure(figsize = (16, 6))
plt.plot(y_harm, color = '#A300F9', label = 'harmonics');
plt.plot(y_perc, color = '#FFB100', label = 'percussive');
plt.legend()
plt.title('Harmonics and percussive for pop')
```

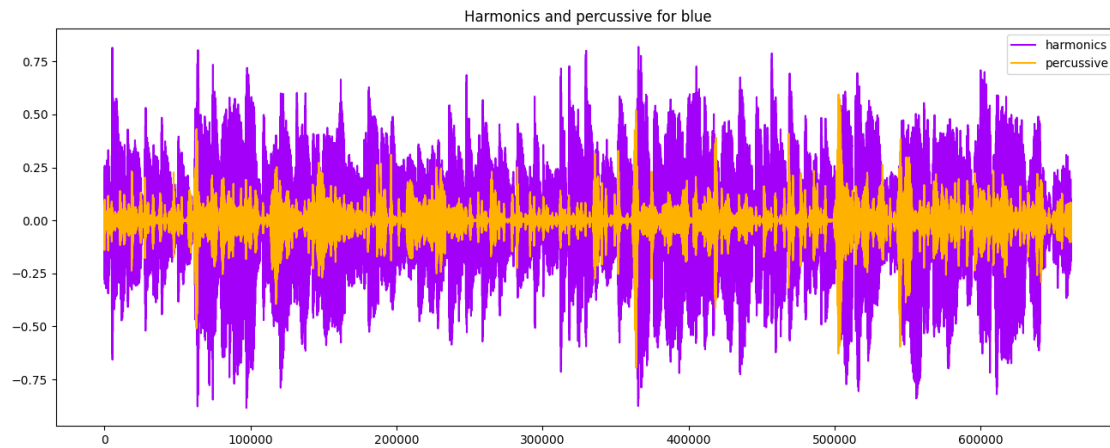
```
[ ]: Text(0.5, 1.0, 'Harmonics and percussive for pop')
```



```
[ ]: y_harm, y_perc = librosa.effects.hpss(y_blue)

plt.figure(figsize = (16, 6))
plt.plot(y_harm, color = '#A300F9', label = 'harmonics');
plt.plot(y_perc, color = '#FFB100', label = 'percussive');
plt.legend()
plt.title('Harmonics and percussive for blue')
```

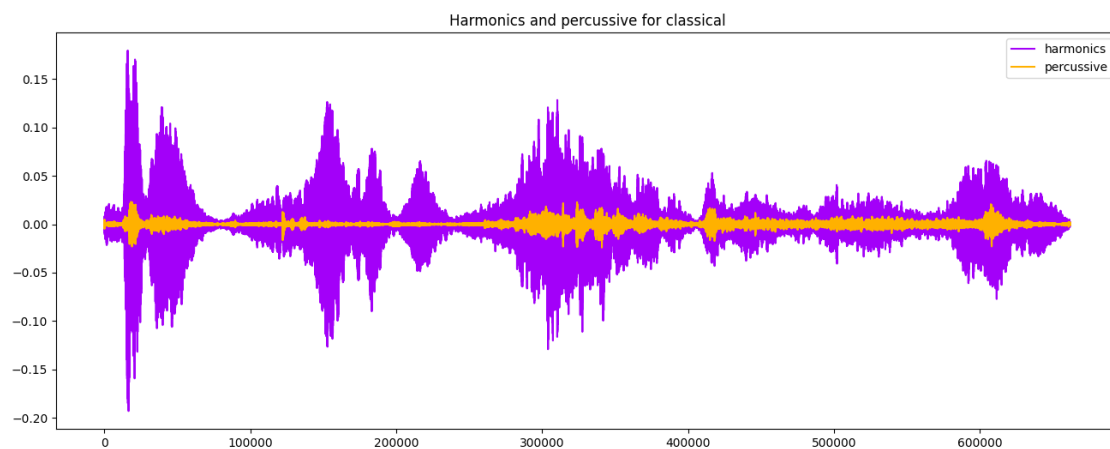
```
[ ]: Text(0.5, 1.0, 'Harmonics and percussive for blue')
```



```
[ ]: y_harm, y_perc = librosa.effects.hpss(y_classical)

plt.figure(figsize = (16, 6))
plt.plot(y_harm, color = '#A300F9', label = 'harmonics');
plt.plot(y_perc, color = '#FFB100', label = 'percussive');
plt.legend()
plt.title('Harmonics and percussive for classical')
```

```
[ ]: Text(0.5, 1.0, 'Harmonics and percussive for classical')
```



### 1.1.7 Tempo BMP (beats per minute)

```
[ ]: tempo, _ = librosa.beat.beat_track(y = y_pop, sr = sr_pop)
tempo
```

```
[ ]: 99.38401442307692
```

### 1.1.8 MFCC

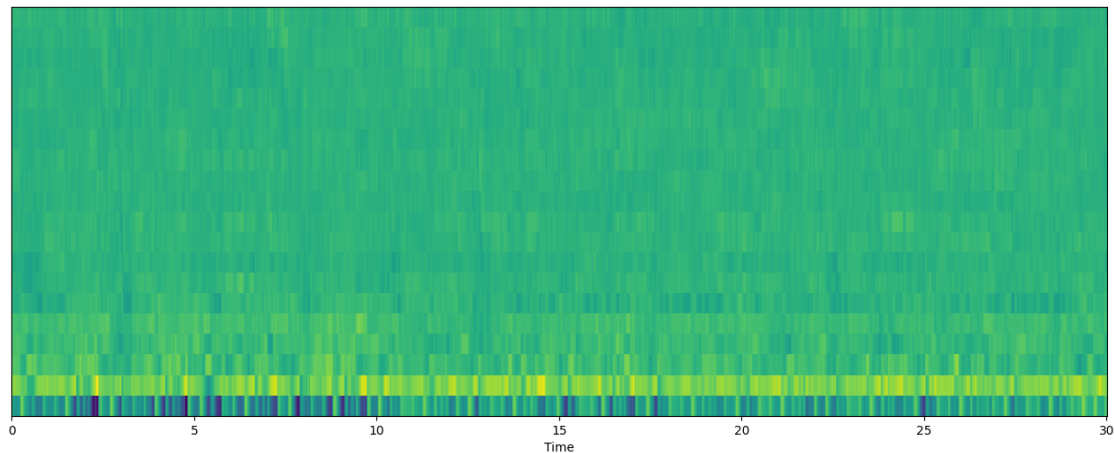
A compressed representation of the Mel Spectrograms by the most essential frequency coef.

The x-axis of the MFCCs plot represents time, and the y-axis represents the MFCC coefficient values.

```
[ ]: mfccs = librosa.feature.mfcc(y = y_pop, sr=sr_pop)
      print('mfccs shape:', mfccs.shape)

      #Displaying the MFCCs:
      plt.figure(figsize = (16, 6))
      librosa.display.specshow(mfccs, sr=sr_pop, x_axis='time', cmap = 'viridis');
```

```
mfccs shape: (20, 1293)
```



```
[ ]: # Perform Feature Scaling
      mfccs = sklearn.preprocessing.scale(mfccs, axis=1)
      print('Mean:', mfccs.mean(), '\n')
      print('Var:', mfccs.var())

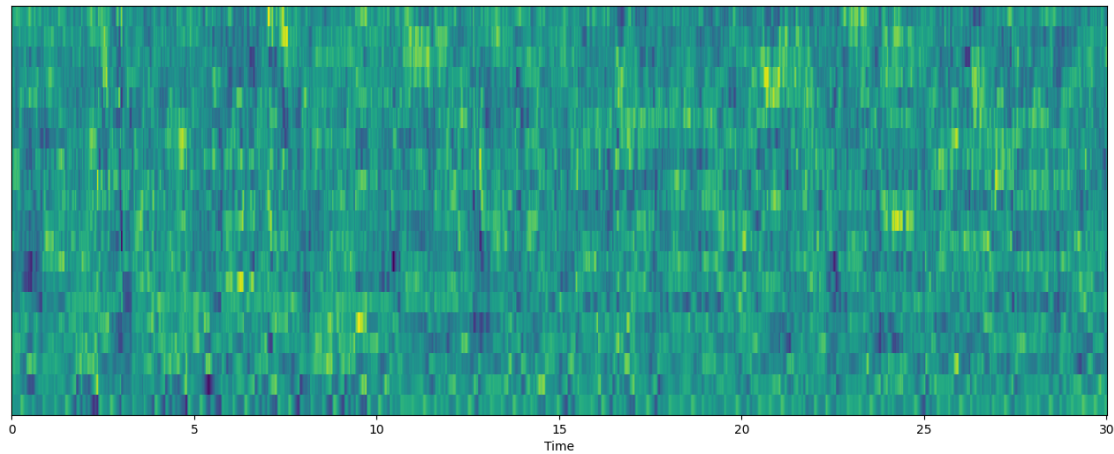
      plt.figure(figsize = (16, 6))
      librosa.display.specshow(mfccs, sr=sr_pop, x_axis='time', cmap = 'viridis');
```

```
Mean: -2.360215e-09
```

```
Var: 1.0
```

```
/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-
packages/sklearn/preprocessing/_data.py:240: UserWarning: Numerical issues were
encountered when centering the data and might not be solved. Dataset may contain
too large values. You may need to prescale your features.
```

```
warnings.warn(
/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-
packages/sklearn/preprocessing/_data.py:259: UserWarning: Numerical issues were
encountered when scaling the data and might not be solved. The standard
deviation of the data is probably very close to 0.
warnings.warn(
```



```
[ ]: print(mfccs)
```

```
[[ 0.09206966  0.38610902  0.11332224 ...  1.2038206   1.370697
   0.82737654]
 [ 0.26141602  0.34526667  0.451189    ...  0.31910515  0.11388601
   0.18868302]
 [ 0.982792    0.8726585   1.0931873   ... -0.43070954 -0.5778949
  -0.76126564]
 ...
 [-0.32814997 -0.04034512 -0.27681282 ...  0.5647969   1.3951571
   2.1499693 ]
 [ 0.73992103  0.7570349   -0.0860804   ... -0.20964907  0.44987693
   1.7667273 ]
 [ 0.89229167  0.74019504  0.55440974 ... -0.8115252  -1.0920751
  -0.89953953]]
```

## 1.2 2. EDA

The dataset contains a list of features from the audio files such as spectral centroid (the weighted mean of the frequencies in an audio file). We use this dataset on machine learning methods to predict classification accuracy later on.

```
[ ]: data = pd.read_csv(f'{path}/features_30_sec.csv')
data.head()
```

```
[ ]:      filename  length  chroma_stft_mean  chroma_stft_var  rms_mean  \
0  blues.00000.wav  661794      0.350088      0.088757  0.130228
1  blues.00001.wav  661794      0.340914      0.094980  0.095948
2  blues.00002.wav  661794      0.363637      0.085275  0.175570
3  blues.00003.wav  661794      0.404785      0.093999  0.141093
4  blues.00004.wav  661794      0.308526      0.087841  0.091529

      rms_var  spectral_centroid_mean  spectral_centroid_var  \
0  0.002827      1784.165850      129774.064525
1  0.002373      1530.176679      375850.073649
2  0.002746      1552.811865      156467.643368
3  0.006346      1070.106615      184355.942417
4  0.002303      1835.004266      343399.939274

      spectral_bandwidth_mean  spectral_bandwidth_var  ...  mfcc16_var  \
0      2002.449060      85882.761315  ...  52.420910
1      2039.036516      213843.755497  ...  55.356403
2      1747.702312      76254.192257  ...  40.598766
3      1596.412872      166441.494769  ...  44.427753
4      1748.172116      88445.209036  ...  86.099236

      mfcc17_mean  mfcc17_var  mfcc18_mean  mfcc18_var  mfcc19_mean  mfcc19_var  \
0      -1.690215   36.524071   -0.408979   41.597103   -2.303523   55.062923
1      -0.731125   60.314529    0.295073   48.120598   -0.283518   51.106190
2      -7.729093   47.639427   -1.816407   52.382141   -3.439720   46.639660
3      -3.319597   50.206673    0.636965   37.319130   -0.619121   37.259739
4      -5.454034   75.269707   -0.916874   53.613918   -4.404827   62.910812

      mfcc20_mean  mfcc20_var  label
0      1.221291   46.936035  blues
1      0.531217   45.786282  blues
2     -2.231258   30.573025  blues
3     -3.407448   31.949339  blues
4    -11.703234   55.195160  blues
```

[5 rows x 60 columns]

### 1.2.1 Correlation heap map

```
[ ]: spike_cols = [col for col in data.columns if 'mean' in col]
      corr = data[spike_cols].corr()

      # Generate a mask for the upper triangle
      mask = np.triu(np.ones_like(corr, dtype=np.bool))

      # Set up the matplotlib figure
      f, ax = plt.subplots(figsize=(16, 11));
```

```

# Generate a custom diverging colormap
cmap = sns.diverging_palette(0, 25, as_cmap=True, s = 90, l = 45, n = 5)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.title('Correlation Heatmap (for the mean variables)', fontsize = 25)
plt.xticks(fontsize = 10)
plt.yticks(fontsize = 10);

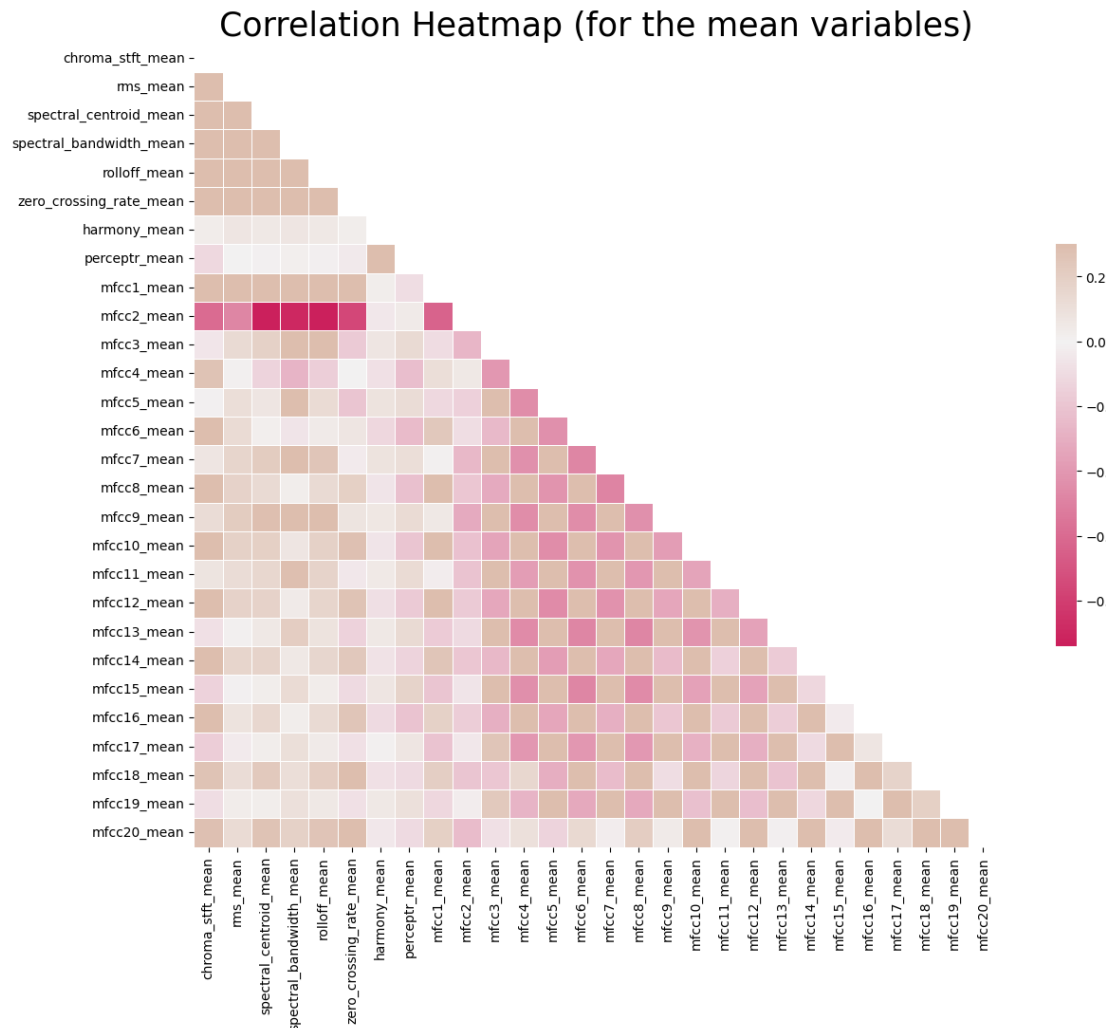
```

/var/folders/0h/xyv81g2n7sj6zr0c9cw30gkc0000gn/T/ipykernel\_81575/4060137719.py:5  
: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To  
silence this warning, use `bool` by itself. Doing this will not modify any  
behavior and is safe. If you specifically wanted the numpy scalar type, use  
`np.bool\_` here.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
mask = np.triu(np.ones_like(corr, dtype=np.bool))
```



It seems that chroma sfft mean, rms mean, spectral centriod mean, spectral bandwidth mean, rolloff mean zero cross rating mean seems to be the most meaningful variables with the strongest correlation relates to each other.

## 1.2.2 PCA

Since there are lots of similarities among features, we use PCA to condense the features and examine the possibilities for classification through componennts.

```
[ ]: from sklearn import preprocessing

data = data.drop(['filename'], axis=1)
y = data['label']
X = data.loc[:, data.columns != 'label']
```


```
#### NORMALIZE X ####
cols = X.columns
min_max_scaler = preprocessing.MinMaxScaler()
np_scaled = min_max_scaler.fit_transform(X)
X = pd.DataFrame(np_scaled, columns = cols)
```

```
[ ]: import numpy as np
      from sklearn.decomposition import PCA
      # Use 100 components to see explained variance

      pca_58 = PCA(n_components = 58)
      pca_58.fit(X)

      plt.grid()
      plt.plot(np.cumsum(pca_58.explained_variance_ratio_ * 100))
      plt.xlabel('Number of components')
      plt.ylabel('Explained variance')
      plt.title('PCA cumulative explained variance')
```

```
[ ]: Text(0.5, 1.0, 'PCA cumulative explained variance')
```





```
[ ]: import pandas as pd
explained_var = np.cumsum(pca_58.explained_variance_ratio_ * 100).reshape(-1, 1)
num_comps = [i+1 for i in range(58)]
explained_var
explained_var_df = pd.DataFrame(explained_var, index=num_comps, columns =
    ↳ ['Explained Variance Ratio'])
explained_var_df
```

```
[ ]:      Explained Variance Ratio
1          24.393550
2          46.175354
3          55.869008
4          62.307090
5          67.594676
6          71.390181
7          74.093713
8          76.294254
9          78.407186
10         80.051094
11         81.390982
12         82.698649
13         83.880329
14         84.954229
15         85.927882
16         86.847171
17         87.707583
18         88.511371
19         89.239098
20         89.931975
21         90.574360
22         91.190001
23         91.759817
24         92.303727
25         92.819800
26         93.305670
27         93.770106
28         94.209590
29         94.638690
30         95.042563
31         95.431485
32         95.789835
33         96.142766
34         96.481475
35         96.805725
36         97.114400
37         97.394488
38         97.642221
```

```

39          97.882790
40          98.111695
41          98.317939
42          98.519024
43          98.706679
44          98.889972
45          99.052849
46          99.198941
47          99.334579
48          99.445599
49          99.536310
50          99.624271
51          99.709509
52          99.787073
53          99.856644
54          99.918138
55          99.958835
56          99.977609
57          99.993447
58          100.000000

```

```

[ ]: from mpl_toolkits import mplot3d
    from sklearn.preprocessing import LabelEncoder

pca = PCA(n_components=3)
principalComponents = pca.fit_transform(X)
principalDf = pd.DataFrame(data = principalComponents, columns = ['principal_
    ↪component 1', 'principal component 2', 'principal component 3'])

# concatenate with target label
finalDf = pd.concat([principalDf, y], axis = 1)

# Define an array of color names
le = LabelEncoder()
color_numbers = le.fit_transform(finalDf["label"])

fig = plt.figure(figsize = (16, 9))
ax = fig.add_subplot(projection='3d')
ax.scatter(finalDf["principal component 1"], finalDf["principal component 2"],
    ↪finalDf["principal component 3"], c = color_numbers, cmap='Paired',
    ↪linewidth=0.5)

# Add labels and title
ax.set_xlabel('Principal Component 1')
ax.set_ylabel('Principal Component 2')
ax.set_zlabel('Principal Component 3')

```

```

ax.set_title('3 Principle component plots')

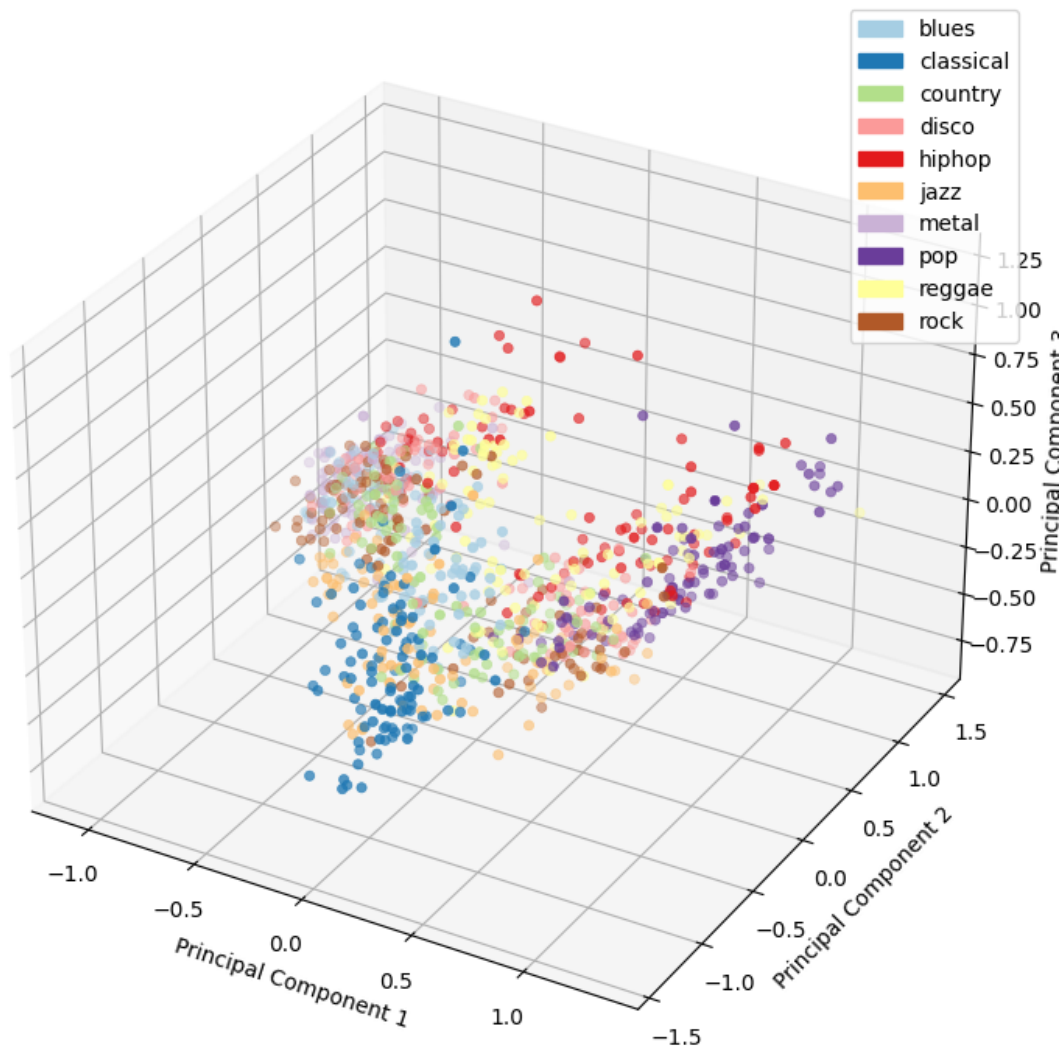
# Find the label names and corresponding color codes
label_names = le.classes_
color_codes = plt.get_cmap('Paired')(np.linspace(0, 1, len(label_names)))

# Create a legend
legend_handles = [plt.Rectangle((0,0),1,1, color=color_codes[i],
    label=label_names[i]) for i in range(len(label_names))]
ax.legend(handles=legend_handles)

plt.show()

```

3 Principle component plots



We can see that there are some patterns but it's not clearly visible to classify different types of music using three components.

### 1.3 3. Classification through feature data

#### 1.3.1 1D CNN classification

We cannot use 2D CNN classifier because of the small amount of features that are extracted from the audio dataset.

```
[ ]: import IPython.display as ipd
import librosa.display

import tensorflow as tf
import tensorflow.keras as keras

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Conv2D

import sklearn.model_selection as sk

from sklearn.model_selection import train_test_split

[ ]: y = data['label'] # genre variable.
X = data.loc[:, data.columns != 'label'] #select all columns but not the labels

#### NORMALIZE X ####
# Normalize so everything is on the same scale.
cols = X.columns
min_max_scaler = preprocessing.MinMaxScaler()
np_scaled = min_max_scaler.fit_transform(X)

# new data frame with the new scaled data.
X = pd.DataFrame(np_scaled, columns = cols)

[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

[ ]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
y_train = label_encoder.fit_transform(y_train)
y_test = label_encoder.fit_transform(y_test)

[ ]: X_train = np.asarray(X_train).reshape((X_train.shape[0], X_train.shape[1], 1))
X_test = np.asarray(X_test).reshape((X_test.shape[0], X_test.shape[1], 1))

# Define the CNN model
```

```

model = keras.Sequential([
    keras.layers.Conv1D(32, kernel_size=3, activation='relu',
        ↪input_shape=X_train.shape[1:]),
    keras.layers.MaxPooling1D(pool_size=2),
    keras.layers.Conv1D(32, kernel_size=3, activation='relu'),
    keras.layers.MaxPooling1D(pool_size=2),
    keras.layers.Flatten(),
    keras.layers.Dense(10, activation='softmax')
])

# Compile the model and train it on the input data
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
    ↪metrics=['accuracy'])
history = model.fit(X_train, y_train, epochs=20, batch_size=32,
    ↪validation_data=(X_test, y_test))

```

Epoch 1/20

2023-04-16 20:38:46.031905: I  
tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113]  
Plugin optimizer for device\_type GPU is enabled.

22/22 [=====] - ETA: 0s - loss: 2.2839 - accuracy:  
0.1200

2023-04-16 20:38:47.090227: I  
tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113]  
Plugin optimizer for device\_type GPU is enabled.

22/22 [=====] - 2s 48ms/step - loss: 2.2839 - accuracy:  
0.1200 - val\_loss: 2.2835 - val\_accuracy: 0.1700

Epoch 2/20

1/22 [>...] - ETA: 0s - loss: 2.2362 - accuracy:  
0.4062Epoch 2/20

22/22 [=====] - 0s 10ms/step - loss: 2.2424 - accuracy:  
0.3100 - val\_loss: 2.2442 - val\_accuracy: 0.2667

Epoch 3/20

22/22 [=====] - 0s 10ms/step - loss: 2.1792 - accuracy:  
0.2700 - val\_loss: 2.1701 - val\_accuracy: 0.2667

Epoch 4/20

22/22 [=====] - 0s 11ms/step - loss: 2.0734 - accuracy:  
0.3414 - val\_loss: 2.0739 - val\_accuracy: 0.3000

Epoch 5/20

22/22 [=====] - 0s 8ms/step - loss: 1.9246 - accuracy:  
0.3886 - val\_loss: 1.9275 - val\_accuracy: 0.3500

Epoch 6/20

22/22 [=====] - 0s 11ms/step - loss: 1.7557 - accuracy:  
0.4286 - val\_loss: 1.8002 - val\_accuracy: 0.3900

Epoch 7/20

```

22/22 [=====] - 0s 8ms/step - loss: 1.6009 - accuracy:
0.4786 - val_loss: 1.6542 - val_accuracy: 0.4833
Epoch 8/20
22/22 [=====] - 0s 8ms/step - loss: 1.4930 - accuracy:
0.4886 - val_loss: 1.5844 - val_accuracy: 0.4700
Epoch 9/20
22/22 [=====] - 0s 9ms/step - loss: 1.4155 - accuracy:
0.5143 - val_loss: 1.5798 - val_accuracy: 0.3767
Epoch 10/20
22/22 [=====] - 0s 9ms/step - loss: 1.3650 - accuracy:
0.5257 - val_loss: 1.4907 - val_accuracy: 0.4667
Epoch 11/20
22/22 [=====] - 0s 8ms/step - loss: 1.2868 - accuracy:
0.5471 - val_loss: 1.4481 - val_accuracy: 0.5000
Epoch 12/20
22/22 [=====] - 0s 8ms/step - loss: 1.2369 - accuracy:
0.5800 - val_loss: 1.4303 - val_accuracy: 0.4833
Epoch 13/20
22/22 [=====] - 0s 8ms/step - loss: 1.1989 - accuracy:
0.5814 - val_loss: 1.3690 - val_accuracy: 0.5167
Epoch 14/20
22/22 [=====] - 0s 8ms/step - loss: 1.1689 - accuracy:
0.6000 - val_loss: 1.3614 - val_accuracy: 0.4967
Epoch 15/20
22/22 [=====] - 0s 10ms/step - loss: 1.1279 - accuracy:
0.6186 - val_loss: 1.3259 - val_accuracy: 0.5200
Epoch 16/20
22/22 [=====] - 0s 8ms/step - loss: 1.1049 - accuracy:
0.6143 - val_loss: 1.3321 - val_accuracy: 0.5200
Epoch 17/20
22/22 [=====] - 0s 9ms/step - loss: 1.0814 - accuracy:
0.6157 - val_loss: 1.3241 - val_accuracy: 0.5267
Epoch 18/20
22/22 [=====] - 0s 9ms/step - loss: 1.0446 - accuracy:
0.6400 - val_loss: 1.2927 - val_accuracy: 0.5367
Epoch 19/20
22/22 [=====] - 0s 10ms/step - loss: 1.0257 - accuracy:
0.6557 - val_loss: 1.2981 - val_accuracy: 0.5467
Epoch 20/20
22/22 [=====] - 0s 8ms/step - loss: 0.9927 - accuracy:
0.6671 - val_loss: 1.2353 - val_accuracy: 0.5700

```

```

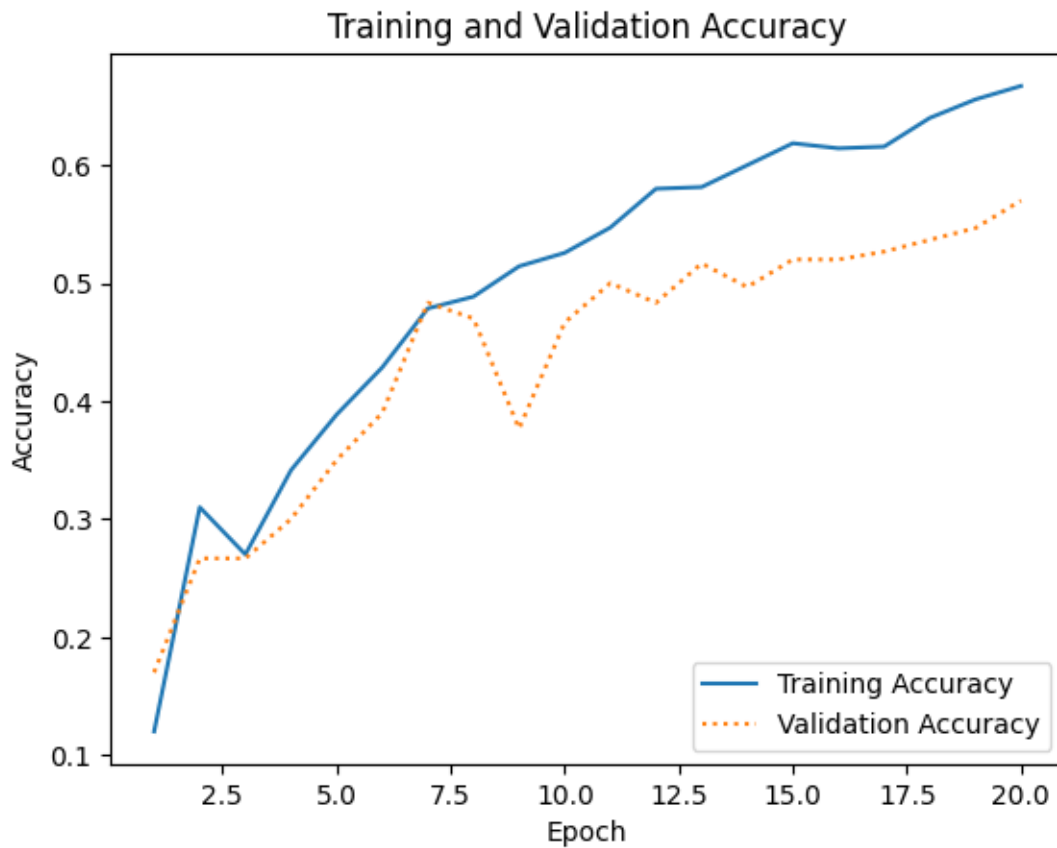
[ ]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, '-', label='Training Accuracy')

```

```
plt.plot(epochs, val_acc, ':', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
```

```
[ ]: <matplotlib.legend.Legend at 0x381033d30>
```



### 1.3.2 Logistic Regression

```
[ ]: X_train = X_train.reshape(X_train.shape[0], X_train.shape[1])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1])

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# Logistic Regression
lg = LogisticRegression(random_state=0, solver='lbfgs',
    ↪multi_class='multinomial')
```

```
lg.fit(X_train, y_train)
y_pred = lg.predict(X_test)

cr = classification_report(y_test, y_pred)
print(cr)
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.46      | 0.46   | 0.46     | 35      |
| 1            | 0.87      | 1.00   | 0.93     | 20      |
| 2            | 0.70      | 0.51   | 0.59     | 37      |
| 3            | 0.65      | 0.44   | 0.53     | 34      |
| 4            | 0.63      | 0.71   | 0.67     | 24      |
| 5            | 0.87      | 0.79   | 0.83     | 33      |
| 6            | 0.53      | 0.90   | 0.67     | 30      |
| 7            | 0.66      | 0.83   | 0.73     | 23      |
| 8            | 0.50      | 0.52   | 0.51     | 29      |
| 9            | 0.48      | 0.34   | 0.40     | 35      |
| accuracy     |           |        | 0.62     | 300     |
| macro avg    | 0.63      | 0.65   | 0.63     | 300     |
| weighted avg | 0.63      | 0.62   | 0.61     | 300     |

/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

### 1.3.3 MLP classifier

```
[ ]: from sklearn.neural_network import MLPClassifier

# Neural Nets
nn = MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(5000, 10),
↳ random_state=1)
nn.fit(X_train, y_train)
y_pred = nn.predict(X_test)

cr = classification_report(y_test, y_pred)
```



```
print(cr)
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.64      | 0.66   | 0.65     | 35      |
| 1            | 0.86      | 0.90   | 0.88     | 20      |
| 2            | 0.43      | 0.24   | 0.31     | 37      |
| 3            | 0.36      | 0.38   | 0.37     | 34      |
| 4            | 0.54      | 0.79   | 0.64     | 24      |
| 5            | 0.81      | 0.67   | 0.73     | 33      |
| 6            | 0.73      | 0.80   | 0.76     | 30      |
| 7            | 0.85      | 0.74   | 0.79     | 23      |
| 8            | 0.48      | 0.48   | 0.48     | 29      |
| 9            | 0.40      | 0.49   | 0.44     | 35      |
| accuracy     |           |        | 0.59     | 300     |
| macro avg    | 0.61      | 0.61   | 0.61     | 300     |
| weighted avg | 0.59      | 0.59   | 0.58     | 300     |

```
/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-  
packages/sklearn/neural_network/_multilayer_perceptron.py:541:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html  
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

### 1.3.4 XGboost classifier

```
[ ]: from xgboost import XGBClassifier, XGBRFClassifier  
from xgboost import plot_tree, plot_importance  
  
# Cross Gradient Booster  
xgb = XGBClassifier(n_estimators=1000, learning_rate=0.05)  
xgb.fit(X_train, y_train)  
y_pred = xgb.predict(X_test)  
  
cr = classification_report(y_test, y_pred)  
print(cr)
```

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.81      | 0.71   | 0.76     | 35      |
| 1 | 0.87      | 1.00   | 0.93     | 20      |
| 2 | 0.79      | 0.70   | 0.74     | 37      |
| 3 | 0.76      | 0.56   | 0.64     | 34      |
| 4 | 0.59      | 0.83   | 0.69     | 24      |

|              |      |      |      |     |
|--------------|------|------|------|-----|
| 5            | 0.93 | 0.82 | 0.87 | 33  |
| 6            | 0.69 | 0.90 | 0.78 | 30  |
| 7            | 0.72 | 0.91 | 0.81 | 23  |
| 8            | 0.68 | 0.59 | 0.63 | 29  |
| 9            | 0.59 | 0.54 | 0.57 | 35  |
| accuracy     |      |      | 0.74 | 300 |
| macro avg    | 0.74 | 0.76 | 0.74 | 300 |
| weighted avg | 0.74 | 0.74 | 0.73 | 300 |

Comparing the accuracy: - XGBoost: 0.74 - Logistic: 0.62 - MLP: 0.59 - 1D CNN: 0.5

We found that XGBoost has the best performance using the feature extracted dataset.

#### 1.4 4. Classification through audio images

```
[ ]: import tensorflow as tf
tf.config.experimental.list_physical_devices('GPU')
```

```
[ ]: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
[ ]: import os
from glob import glob
from PIL import Image
import cv2
import os
from sklearn.model_selection import train_test_split
from sklearn.svm import LinearSVC
import tensorflow as tf
from tensorflow import keras
from sklearn.metrics import classification_report
# from resizeimage import resizeimage
import numpy as np
import matplotlib.pyplot as plt
```

```
[ ]: input_dir = 'GTZAN/images_original/'

# Define input directory
def process_images(path):
    images = []
    flattened = []
    # Loop over all images in input directory
    for filename in os.listdir(path):
        # Load image
        image_path = os.path.join(path, filename)
        image = cv2.imread(image_path)
        # print(image.shape)
```

```

        resize = cv2.resize(image, (100, 70))
        images.append(resize)
        flattened.append(np.array(resize).flatten())
    X = np.asarray(flattened)
    return images, flattened, X

def process_categories(input_dir):
    X_categories = []
    y_labels = []

    for i,subdir in enumerate(os.listdir(input_dir)):
        images, flattened, X = process_images(os.path.join(input_dir, subdir))
        X_categories.append(X)
        y_labels.append(np.full(len(X), i))
    X = np.vstack(X_categories)
    y = np.concatenate(y_labels)
    return X, y

X, y = process_categories(input_dir)

```

```

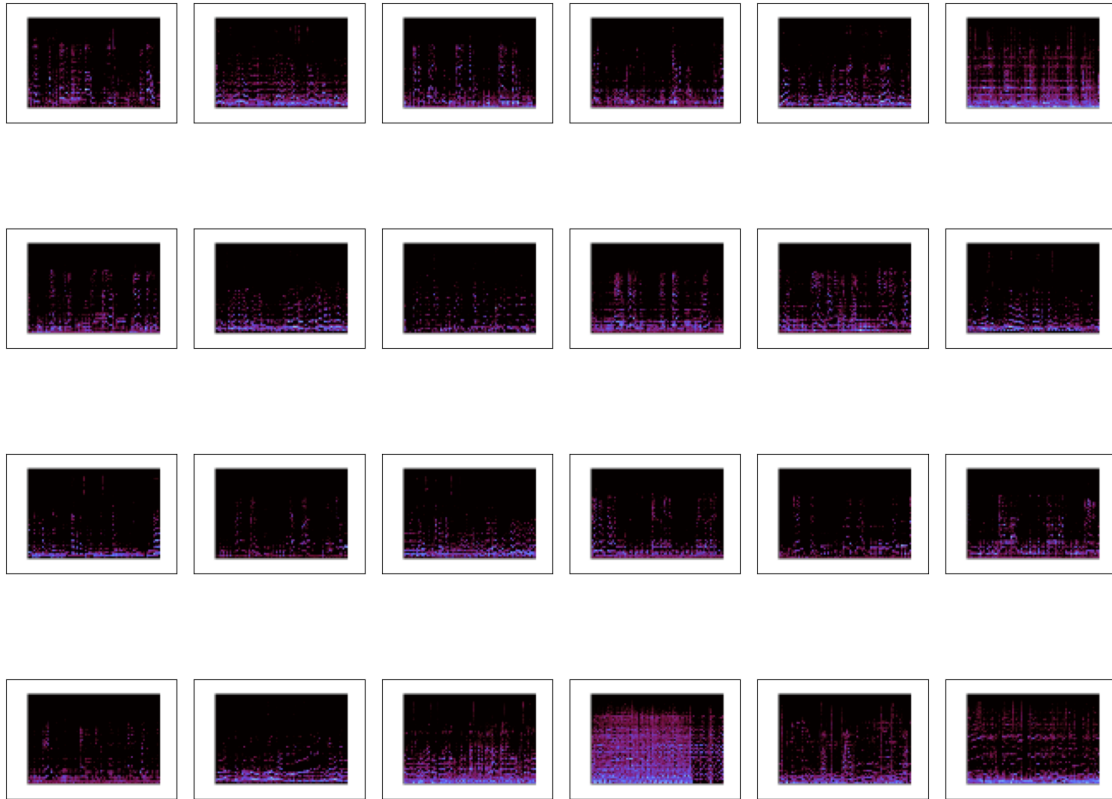
[ ]: blue_dir = 'GTZAN/images_original/blues'
images_blue, flattened_blue , X_blue = process_images(blue_dir)

#plot the photos
def original_photos(images):
    fig, ax = plt.subplots(4, 6, figsize=(20, 16),
                           subplot_kw={'xticks':[], 'yticks':[]},
                           gridspec_kw=dict(hspace=0.1, wspace=0.1))

    k = 0
    for i in range(4):
        for j in range(6):
            ax[i, j].imshow(images[k])
            k += 1

    plt.show()
original_photos(images_blue)

```



```
[ ]: #train test split
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.20,
↳random_state=123)
print(f'There are {len(X_train)} training images, and {len(X_test)} testing_
↳images.' )
```

There are 799 training images, and 200 testing images.

## 1.5 Logistic Regression: baseline analysis

```
[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# Logistic Regression
lg = LogisticRegression(random_state=0, solver='lbfgs',
    ↪multi_class='multinomial')

lg.fit(X_train, y_train)
y_pred = lg.predict(X_test)

cr = classification_report(y_test, y_pred)
print(cr)
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.39      | 0.67   | 0.49     | 18      |
| 1            | 0.79      | 0.62   | 0.70     | 24      |
| 2            | 0.24      | 0.24   | 0.24     | 17      |
| 3            | 0.58      | 0.54   | 0.56     | 28      |
| 4            | 0.30      | 0.46   | 0.36     | 13      |
| 5            | 0.89      | 0.53   | 0.67     | 15      |
| 6            | 0.27      | 0.39   | 0.32     | 18      |
| 7            | 0.54      | 0.35   | 0.42     | 20      |
| 8            | 0.41      | 0.46   | 0.43     | 24      |
| 9            | 0.75      | 0.39   | 0.51     | 23      |
| accuracy     |           |        | 0.47     | 200     |
| macro avg    | 0.51      | 0.46   | 0.47     | 200     |
| weighted avg | 0.53      | 0.47   | 0.48     | 200     |

/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

## 1.6 MLP Classifier

MLP classifier doesn't able to run.

```
[ ]: # from sklearn.neural_network import MLPClassifier

# # Neural Nets
# nn = MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(5000, 10),
    ↪ random_state=1)
# nn.fit(X_train, y_train)
# y_pred = nn.predict(X_test)

# cr = classification_report(y_test, y_pred)
# print(cr)
```

## 1.7 Linear SVC

```
[ ]: svc = LinearSVC(verbose=0).fit(X_train, y_train)

score = svc.score(X_train, y_train)
print("Score: ", score)

#prediction
y_pred = svc.predict(X_test)

cr = classification_report(y_test, y_pred)
print(cr)
```

Score: 0.9987484355444305

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.42      | 0.56   | 0.48     | 18      |
| 1            | 0.72      | 0.75   | 0.73     | 24      |
| 2            | 0.25      | 0.24   | 0.24     | 17      |
| 3            | 0.47      | 0.32   | 0.38     | 28      |
| 4            | 0.25      | 0.31   | 0.28     | 13      |
| 5            | 0.43      | 0.60   | 0.50     | 15      |
| 6            | 0.25      | 0.28   | 0.26     | 18      |
| 7            | 0.36      | 0.25   | 0.29     | 20      |
| 8            | 0.36      | 0.38   | 0.37     | 24      |
| 9            | 0.50      | 0.43   | 0.47     | 23      |
| accuracy     |           |        |          | 0.41    |
| macro avg    |           |        |          | 0.41    |
| weighted avg |           |        |          | 0.41    |

```
/Users/swimmingcircle/miniforge3/envs/tf/lib/python3.9/site-
packages/sklearn/svm/_base.py:1244: ConvergenceWarning: Liblinear failed to
converge, increase the number of iterations.
warnings.warn(
```

## 1.8 CNN

```
[ ]: X_train.shape
```

```
[ ]: (799, 21000)
```

```
[ ]: import numpy as np
X_train_resaped = np.reshape(X_train, (X_train.shape[0], 100, 70, 3))
X_test_resaped = np.reshape(X_test, (X_test.shape[0], 100, 70, 3))

# X_train_resaped = np.reshape(X_train, (X_train.shape[0], 432, 288, 3))
# X_test_resaped = np.reshape(X_test, (X_test.shape[0], 432, 288, 3))
X_train_resaped.shape, X_test_resaped.shape
```

```
[ ]: ((799, 100, 70, 3), (200, 100, 70, 3))
```

```
[ ]: import numpy as np

# Reshape x_train to a 4D tensor
height = 100
width = 70
channels = 3

# Convert x_train to float32 data type and normalize the pixel values to [0, 1]
X_train_resaped = X_train_resaped.astype(np.float32) / 255.0
X_test_resaped = X_test_resaped.astype(np.float32) / 255.0
y_train_resaped = y_train.reshape(-1,1)
y_test_resaped = y_test.reshape(-1,1)
```

```
[ ]: import tensorflow as tf

# Define input shape
input_shape = (height, width, channels)

# Define model
model1 = keras.Sequential([
    keras.layers.Conv2D(32, kernel_size=(3, 3), activation='relu',
↳ input_shape=input_shape),
    keras.layers.MaxPooling2D(pool_size=(2, 2)),
    keras.layers.Flatten(),
    keras.layers.Dense(216, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(10, activation='softmax')
])

# Compile model
model1.compile(optimizer='adam', loss='categorical_crossentropy',
↳ metrics=['accuracy'])
```

```
# Print model summary
model1.summary()
```

Model: "sequential\_36"

| Layer (type)                    | Output Shape       | Param #  |
|---------------------------------|--------------------|----------|
| conv2d_60 (Conv2D)              | (None, 98, 68, 32) | 896      |
| max_pooling2d_54 (MaxPooling2D) | (None, 49, 34, 32) | 0        |
| flatten_37 (Flatten)            | (None, 53312)      | 0        |
| dense_92 (Dense)                | (None, 216)        | 11515608 |
| dropout_56 (Dropout)            | (None, 216)        | 0        |
| dense_93 (Dense)                | (None, 10)         | 2170     |

```
=====  
Total params: 11,518,674  
Trainable params: 11,518,674  
Non-trainable params: 0  
=====
```

| Layer (type)                    | Output Shape       | Param #  |
|---------------------------------|--------------------|----------|
| conv2d_60 (Conv2D)              | (None, 98, 68, 32) | 896      |
| max_pooling2d_54 (MaxPooling2D) | (None, 49, 34, 32) | 0        |
| flatten_37 (Flatten)            | (None, 53312)      | 0        |
| dense_92 (Dense)                | (None, 216)        | 11515608 |
| dropout_56 (Dropout)            | (None, 216)        | 0        |
| dense_93 (Dense)                | (None, 10)         | 2170     |

```
=====  
Total params: 11,518,674  
Trainable params: 11,518,674  
Non-trainable params: 0  
=====
```



```
[ ]: from keras.utils import to_categorical
```

```
y_train_one_hot = to_categorical(y_train_reshaped)  
y_test_one_hot = to_categorical(y_test_reshaped)
```

```
[ ]: history1 = model1.fit(X_train_reshaped, y_train_one_hot, epochs=25,  
    ↪ validation_data=(X_test_reshaped, y_test_one_hot))
```

Epoch 1/25

2023-04-16 19:45:03.906596: I  
tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113]  
Plugin optimizer for device\_type GPU is enabled.

25/25 [=====] - ETA: 0s - loss: 4.9602 - accuracy:  
0.0989

2023-04-16 19:45:05.676026: I  
tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113]  
Plugin optimizer for device\_type GPU is enabled.

25/25 [=====] - 2s 71ms/step - loss: 4.9602 - accuracy:  
0.0989 - val\_loss: 2.3149 - val\_accuracy: 0.0900

Epoch 2/25

Epoch 2/25

25/25 [=====] - 1s 44ms/step - loss: 2.2831 - accuracy:  
0.1089 - val\_loss: 2.2716 - val\_accuracy: 0.1550

Epoch 3/25

25/25 [=====] - 1s 43ms/step - loss: 2.2654 - accuracy:  
0.1277 - val\_loss: 2.2515 - val\_accuracy: 0.1200

Epoch 4/25

25/25 [=====] - 1s 42ms/step - loss: 2.2073 - accuracy:  
0.1252 - val\_loss: 2.2227 - val\_accuracy: 0.1200

Epoch 5/25

25/25 [=====] - 1s 42ms/step - loss: 2.1815 - accuracy:  
0.1827 - val\_loss: 2.1425 - val\_accuracy: 0.1800

Epoch 6/25

25/25 [=====] - 1s 44ms/step - loss: 2.1531 - accuracy:  
0.2028 - val\_loss: 2.1401 - val\_accuracy: 0.2050

Epoch 7/25

25/25 [=====] - 1s 44ms/step - loss: 2.1339 - accuracy:  
0.1815 - val\_loss: 2.0889 - val\_accuracy: 0.2450

Epoch 8/25

25/25 [=====] - 1s 41ms/step - loss: 2.0557 - accuracy:  
0.1865 - val\_loss: 2.0710 - val\_accuracy: 0.1900

Epoch 9/25

25/25 [=====] - 1s 43ms/step - loss: 2.0235 - accuracy:  
0.2491 - val\_loss: 2.0074 - val\_accuracy: 0.2550

Epoch 10/25


25/25 [=====] - 1s 41ms/step - loss: 1.9714 - accuracy:

0.2703 - val\_loss: 1.9570 - val\_accuracy: 0.2850  
Epoch 11/25  
25/25 [=====] - 1s 42ms/step - loss: 1.8889 - accuracy:  
0.2904 - val\_loss: 1.9128 - val\_accuracy: 0.2350  
Epoch 12/25  
25/25 [=====] - 1s 42ms/step - loss: 1.8125 - accuracy:  
0.3191 - val\_loss: 1.8602 - val\_accuracy: 0.2850  
Epoch 13/25  
25/25 [=====] - 1s 42ms/step - loss: 1.7875 - accuracy:  
0.3379 - val\_loss: 1.8296 - val\_accuracy: 0.2950  
Epoch 14/25  
25/25 [=====] - 1s 41ms/step - loss: 1.7034 - accuracy:  
0.3792 - val\_loss: 1.9135 - val\_accuracy: 0.2650  
Epoch 15/25  
25/25 [=====] - 1s 42ms/step - loss: 1.6393 - accuracy:  
0.3892 - val\_loss: 1.7601 - val\_accuracy: 0.3350  
Epoch 16/25  
25/25 [=====] - 1s 41ms/step - loss: 1.5424 - accuracy:  
0.4618 - val\_loss: 1.7579 - val\_accuracy: 0.3100  
Epoch 17/25  
25/25 [=====] - 1s 42ms/step - loss: 1.5447 - accuracy:  
0.4481 - val\_loss: 1.6982 - val\_accuracy: 0.3150  
Epoch 18/25  
25/25 [=====] - 1s 41ms/step - loss: 1.3669 - accuracy:  
0.5181 - val\_loss: 1.6034 - val\_accuracy: 0.4450  
Epoch 19/25  
25/25 [=====] - 1s 42ms/step - loss: 1.3153 - accuracy:  
0.5419 - val\_loss: 1.6054 - val\_accuracy: 0.4050  
Epoch 20/25  
25/25 [=====] - 1s 43ms/step - loss: 1.2225 - accuracy:  
0.5982 - val\_loss: 1.6541 - val\_accuracy: 0.3850  
Epoch 21/25  
25/25 [=====] - 1s 42ms/step - loss: 1.1374 - accuracy:  
0.6108 - val\_loss: 1.5364 - val\_accuracy: 0.4500  
Epoch 22/25  
25/25 [=====] - 1s 42ms/step - loss: 1.0269 - accuracy:  
0.6471 - val\_loss: 1.5674 - val\_accuracy: 0.3950  
Epoch 23/25  
25/25 [=====] - 2s 70ms/step - loss: 0.9930 - accuracy:  
0.6646 - val\_loss: 1.5558 - val\_accuracy: 0.4700  
Epoch 24/25  
25/25 [=====] - 1s 44ms/step - loss: 0.9360 - accuracy:  
0.6934 - val\_loss: 1.4904 - val\_accuracy: 0.4600  
Epoch 25/25  
25/25 [=====] - 1s 42ms/step - loss: 0.8280 - accuracy:  
0.7447 - val\_loss: 1.5228 - val\_accuracy: 0.4050

```
[ ]: acc = history1.history['accuracy']
val_acc = history1.history['val_accuracy']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, '-', label='Training Accuracy')
plt.plot(epochs, val_acc, ':', label='Validation Accuracy')
plt.title('Training and Validation Accuracy for model 1')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
```

```
[ ]: <matplotlib.legend.Legend at 0x2db138a90>
```



```
[ ]: from keras.layers import LeakyReLU

# Define model
reg_param = 0.001

model2 = keras.Sequential([
```

```

keras.layers.Conv2D(64, kernel_size=(3, 3), input_shape=input_shape),
LeakyReLU(alpha=0.1),
keras.layers.MaxPooling2D(pool_size=(2, 2)),
keras.layers.Flatten(),
keras.layers.Dense(216),
LeakyReLU(alpha=0.1),
keras.layers.Dropout(0.2),
keras.layers.Dense(128),
LeakyReLU(alpha=0.1),
keras.layers.Dropout(0.1),
keras.layers.Dense(10, activation='softmax')
])

# Compile model
model2.compile(optimizer='adam', loss='categorical_crossentropy',
               metrics=['accuracy'])

# Print model summary
model2.summary()

```

Model: "sequential\_30"

| Layer (type)                    | Output Shape       | Param #  |
|---------------------------------|--------------------|----------|
| conv2d_42 (Conv2D)              | (None, 98, 68, 64) | 1792     |
| leaky_re_lu_7 (LeakyReLU)       | (None, 98, 68, 64) | 0        |
| max_pooling2d_42 (MaxPooling2D) | (None, 49, 34, 64) | 0        |
| flatten_31 (Flatten)            | (None, 106624)     | 0        |
| dense_76 (Dense)                | (None, 216)        | 23031000 |
| leaky_re_lu_8 (LeakyReLU)       | (None, 216)        | 0        |
| dropout_46 (Dropout)            | (None, 216)        | 0        |
| dense_77 (Dense)                | (None, 128)        | 27776    |
| leaky_re_lu_9 (LeakyReLU)       | (None, 128)        | 0        |
| dropout_47 (Dropout)            | (None, 128)        | 0        |
| conv2d_42 (Conv2D)              | (None, 98, 68, 64) | 1792     |

|                                 |                    |          |
|---------------------------------|--------------------|----------|
| leaky_re_lu_7 (LeakyReLU)       | (None, 98, 68, 64) | 0        |
| max_pooling2d_42 (MaxPooling2D) | (None, 49, 34, 64) | 0        |
| flatten_31 (Flatten)            | (None, 106624)     | 0        |
| dense_76 (Dense)                | (None, 216)        | 23031000 |
| leaky_re_lu_8 (LeakyReLU)       | (None, 216)        | 0        |
| dropout_46 (Dropout)            | (None, 216)        | 0        |
| dense_77 (Dense)                | (None, 128)        | 27776    |
| leaky_re_lu_9 (LeakyReLU)       | (None, 128)        | 0        |
| dropout_47 (Dropout)            | (None, 128)        | 0        |
| dense_78 (Dense)                | (None, 10)         | 1290     |

```
=====
Total params: 23,061,858
Trainable params: 23,061,858
Non-trainable params: 0
-----
```

```
[ ]: history2 = model2.fit(X_train_reshaped, y_train_one_hot, epochs=20,
    ↪ validation_data=(X_test_reshaped, y_test_one_hot))
```

Epoch 1/20

```
2023-04-16 19:38:25.937693: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
25/25 [=====] - ETA: 0s - loss: 5.2792 - accuracy:
0.1239
```

```
2023-04-16 19:38:28.773806: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
25/25 [=====] - 3s 113ms/step - loss: 5.2792 -
accuracy: 0.1239 - val_loss: 2.1029 - val_accuracy: 0.2500
```

Epoch 2/20

Epoch 2/20

```
25/25 [=====] - 2s 96ms/step - loss: 2.1115 - accuracy:
0.2403 - val_loss: 2.0683 - val_accuracy: 0.1900
```

Epoch 3/20  
25/25 [=====] - 2s 89ms/step - loss: 1.9993 - accuracy: 0.2628 - val\_loss: 1.9821 - val\_accuracy: 0.2550

Epoch 4/20  
25/25 [=====] - 2s 89ms/step - loss: 1.8890 - accuracy: 0.3041 - val\_loss: 1.9971 - val\_accuracy: 0.2250

Epoch 5/20  
25/25 [=====] - 2s 87ms/step - loss: 1.7691 - accuracy: 0.3479 - val\_loss: 1.8451 - val\_accuracy: 0.2850

Epoch 6/20  
25/25 [=====] - 2s 87ms/step - loss: 1.6079 - accuracy: 0.4543 - val\_loss: 1.7730 - val\_accuracy: 0.3350

Epoch 7/20  
25/25 [=====] - 2s 88ms/step - loss: 1.3617 - accuracy: 0.5244 - val\_loss: 1.6710 - val\_accuracy: 0.3700

Epoch 8/20  
25/25 [=====] - 2s 87ms/step - loss: 1.0907 - accuracy: 0.6258 - val\_loss: 1.5692 - val\_accuracy: 0.4050

Epoch 9/20  
25/25 [=====] - 2s 96ms/step - loss: 0.8184 - accuracy: 0.7322 - val\_loss: 1.4696 - val\_accuracy: 0.4500

Epoch 10/20  
25/25 [=====] - 2s 90ms/step - loss: 0.6414 - accuracy: 0.7897 - val\_loss: 1.6548 - val\_accuracy: 0.4050

Epoch 11/20  
25/25 [=====] - 2s 90ms/step - loss: 0.5033 - accuracy: 0.8360 - val\_loss: 1.4528 - val\_accuracy: 0.4850

Epoch 12/20  
25/25 [=====] - 2s 87ms/step - loss: 0.2672 - accuracy: 0.9387 - val\_loss: 1.4447 - val\_accuracy: 0.4650

Epoch 13/20  
25/25 [=====] - 2s 89ms/step - loss: 0.1924 - accuracy: 0.9449 - val\_loss: 1.4954 - val\_accuracy: 0.5100

Epoch 14/20  
25/25 [=====] - 2s 87ms/step - loss: 0.1153 - accuracy: 0.9787 - val\_loss: 1.4490 - val\_accuracy: 0.5400

Epoch 15/20  
25/25 [=====] - 2s 96ms/step - loss: 0.0741 - accuracy: 0.9875 - val\_loss: 1.5289 - val\_accuracy: 0.4950

Epoch 16/20  
25/25 [=====] - 2s 92ms/step - loss: 0.0448 - accuracy: 0.9975 - val\_loss: 1.4922 - val\_accuracy: 0.5050

Epoch 17/20  
25/25 [=====] - 2s 89ms/step - loss: 0.0477 - accuracy: 0.9912 - val\_loss: 1.8398 - val\_accuracy: 0.4750

Epoch 18/20  
25/25 [=====] - 2s 95ms/step - loss: 0.0356 - accuracy: 0.9950 - val\_loss: 1.4675 - val\_accuracy: 0.5400

Epoch 19/20

25/25 [=====] - 2s 89ms/step - loss: 0.0250 - accuracy: 0.9975 - val\_loss: 1.5995 - val\_accuracy: 0.5150

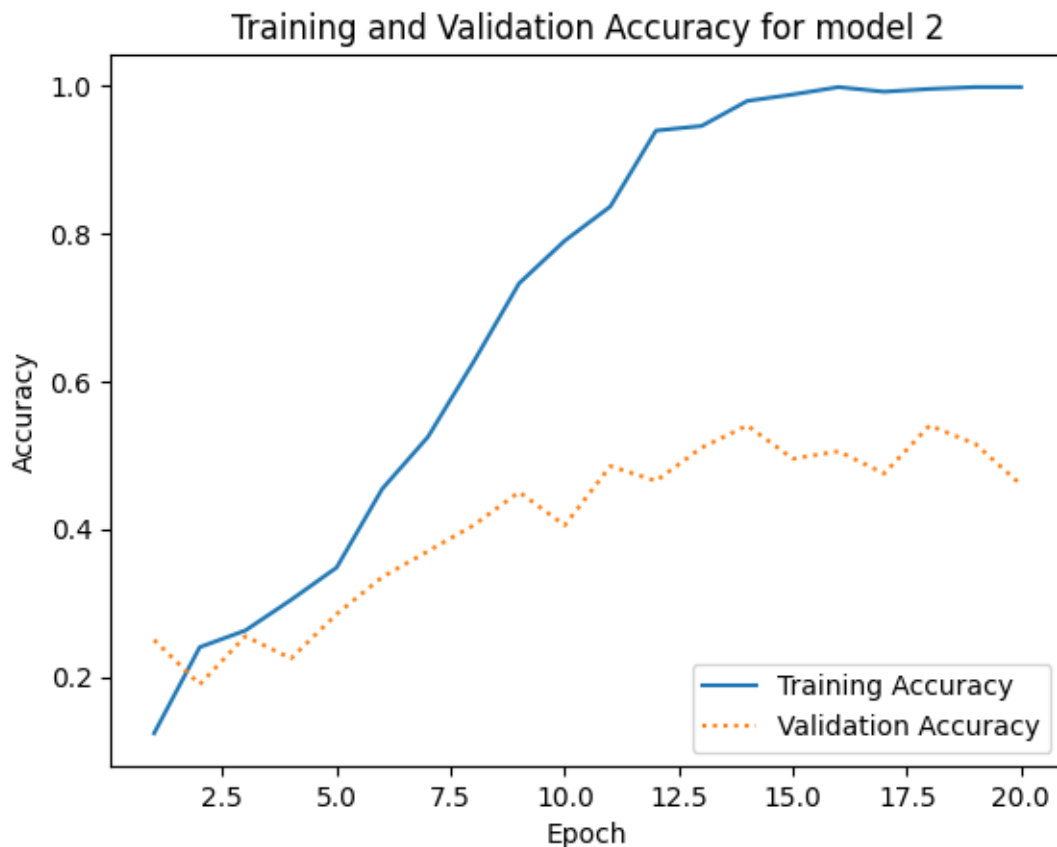
Epoch 20/20

25/25 [=====] - 2s 87ms/step - loss: 0.0175 - accuracy: 0.9975 - val\_loss: 1.7582 - val\_accuracy: 0.4600

```
[ ]: acc = history2.history['accuracy']
val_acc = history2.history['val_accuracy']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, '-', label='Training Accuracy')
plt.plot(epochs, val_acc, ':', label='Validation Accuracy')
plt.title('Training and Validation Accuracy for model 2')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
```

[ ]: <matplotlib.legend.Legend at 0x2cdd466a0>



```
[ ]: import tensorflow as tf

# Define the AlexNet model
def AlexNet(input_shape, num_classes):
    model = tf.keras.models.Sequential([
        # Convolutional layer 1
        tf.keras.layers.Conv2D(filters=96, kernel_size=(11,11), strides=(4,4),
        ↪padding='valid', activation='relu', input_shape=input_shape),
        tf.keras.layers.MaxPooling2D(pool_size=(3,3), strides=(2,2),
        ↪padding='valid'),
        tf.keras.layers.BatchNormalization(),

        # Convolutional layer 2
        tf.keras.layers.Conv2D(filters=256, kernel_size=(5,5), strides=(1,1),
        ↪padding='same', activation='relu'),
        tf.keras.layers.MaxPooling2D(pool_size=(3,3), strides=(2,2),
        ↪padding='valid'),
        tf.keras.layers.BatchNormalization(),

        # Convolutional layer 3
        tf.keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1),
        ↪padding='same', activation='relu'),
        tf.keras.layers.BatchNormalization(),

        # Convolutional layer 4
        tf.keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1),
        ↪padding='same', activation='relu'),
        tf.keras.layers.BatchNormalization(),

        # Convolutional layer 5
        tf.keras.layers.Conv2D(filters=256, kernel_size=(3,3), strides=(1,1),
        ↪padding='same', activation='relu'),
        tf.keras.layers.MaxPooling2D(pool_size=(3,3), strides=(2,2),
        ↪padding='valid'),
        tf.keras.layers.BatchNormalization(),

        # Flatten the output from the convolutional layers
        tf.keras.layers.Flatten(),

        # Fully connected layer 1
        tf.keras.layers.Dense(units=4096, activation='relu'),
        tf.keras.layers.Dropout(0.5),

        # Fully connected layer 2
        tf.keras.layers.Dense(units=4096, activation='relu'),
        tf.keras.layers.Dropout(0.5),
```



```

        # Output layer
        tf.keras.layers.Dense(units=num_classes, activation='softmax')
    ])

    return model
model3 = AlexNet(input_shape, 10)
# Compile model
model3.compile(optimizer='adam', loss='categorical_crossentropy',
               metrics=['accuracy'])

# Print model summary
model3.summary()

```

Model: "sequential\_33"

| Layer (type)                                 | Output Shape       | Param # |
|--|--------------------|---------|
| conv2d_53 (Conv2D)                           | (None, 23, 15, 96) | 34944   |
| max_pooling2d_49 (MaxPooling2D)              | (None, 11, 7, 96)  | 0       |
| batch_normalization_21 (Batch Normalization) | (None, 11, 7, 96)  | 384     |
| conv2d_54 (Conv2D)                           | (None, 11, 7, 256) | 614656  |
| max_pooling2d_50 (MaxPooling2D)              | (None, 5, 3, 256)  | 0       |
| batch_normalization_22 (Batch Normalization) | (None, 5, 3, 256)  | 1024    |
| conv2d_55 (Conv2D)                           | (None, 5, 3, 384)  | 885120  |
| batch_normalization_23 (Batch Normalization) | (None, 5, 3, 384)  | 1536    |
| conv2d_56 (Conv2D)                           | (None, 5, 3, 384)  | 1327488 |
| batch_normalization_24 (Batch Normalization) | (None, 5, 3, 384)  | 1536    |
| conv2d_57 (Conv2D)                           | (None, 5, 3, 256)  | 884992  |
| max_pooling2d_51 (MaxPooling2D)              | (None, 2, 1, 256)  | 0       |

|  |                    |         |
|--|--------------------|---------|
| batch_normalization_25 (Batch Normalization) | (None, 2, 1, 256)  | 1024    |
| flatten_34 (Flatten)                         | (None, 512)        | 0       |
| dense_85 (Dense)                             | (None, 4096)       | 2101248 |
| -----  |                    |         |
| Layer (type)                                 | Output Shape       | Param # |
| =====  |                    |         |
| conv2d_53 (Conv2D)                           | (None, 23, 15, 96) | 34944   |
| max_pooling2d_49 (MaxPooling2D)              | (None, 11, 7, 96)  | 0       |
| batch_normalization_21 (Batch Normalization) | (None, 11, 7, 96)  | 384     |
| conv2d_54 (Conv2D)                           | (None, 11, 7, 256) | 614656  |
| max_pooling2d_50 (MaxPooling2D)              | (None, 5, 3, 256)  | 0       |
| batch_normalization_22 (Batch Normalization) | (None, 5, 3, 256)  | 1024    |
| conv2d_55 (Conv2D)                           | (None, 5, 3, 384)  | 885120  |
| batch_normalization_23 (Batch Normalization) | (None, 5, 3, 384)  | 1536    |
| conv2d_56 (Conv2D)                           | (None, 5, 3, 384)  | 1327488 |
| batch_normalization_24 (Batch Normalization) | (None, 5, 3, 384)  | 1536    |
| conv2d_57 (Conv2D)                           | (None, 5, 3, 256)  | 884992  |
| max_pooling2d_51 (MaxPooling2D)              | (None, 2, 1, 256)  | 0       |
| batch_normalization_25 (Batch Normalization) | (None, 2, 1, 256)  | 1024    |
| flatten_34 (Flatten)                         | (None, 512)        | 0       |
| dense_85 (Dense)                             | (None, 4096)       | 2101248 |

|                      |              |          |
|----------------------|--------------|----------|
| dropout_52 (Dropout) | (None, 4096) | 0        |
| dense_86 (Dense)     | (None, 4096) | 16781312 |
| dropout_53 (Dropout) | (None, 4096) | 0        |
| dense_87 (Dense)     | (None, 10)   | 40970    |

```
=====
Total params: 22,676,234
Trainable params: 22,673,482
Non-trainable params: 2,752
-----
```

```
[ ]: history3 = model3.fit(X_train_resaped, y_train_one_hot, epochs=20,
    ↪validation_data=(X_test_resaped, y_test_one_hot))
```

Epoch 1/20

```
2023-04-16 19:39:37.282980: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
25/25 [=====] - ETA: 0s - loss: 3.0315 - accuracy:
0.1802
```

```
2023-04-16 19:39:43.096270: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.
```

```
25/25 [=====] - 7s 146ms/step - loss: 3.0315 -
accuracy: 0.1802 - val_loss: 5.9092 - val_accuracy: 0.0650
```

Epoch 2/20

Epoch 2/20

```
25/25 [=====] - 2s 68ms/step - loss: 2.1320 - accuracy:
0.2365 - val_loss: 4.2810 - val_accuracy: 0.0650
```

Epoch 3/20

```
25/25 [=====] - 2s 68ms/step - loss: 1.9450 - accuracy:
0.2879 - val_loss: 4.4332 - val_accuracy: 0.0750
```

Epoch 4/20

```
25/25 [=====] - 2s 64ms/step - loss: 1.8114 - accuracy:
0.3229 - val_loss: 4.7994 - val_accuracy: 0.0900
```

Epoch 5/20

```
25/25 [=====] - 2s 62ms/step - loss: 1.7209 - accuracy:
0.3680 - val_loss: 10.2604 - val_accuracy: 0.0750
```

Epoch 6/20

```
25/25 [=====] - 2s 67ms/step - loss: 1.6505 - accuracy:
0.3917 - val_loss: 6.0424 - val_accuracy: 0.1800
```

Epoch 7/20

```

25/25 [=====] - 2s 64ms/step - loss: 1.5404 - accuracy:
0.4293 - val_loss: 5.5076 - val_accuracy: 0.1550
Epoch 8/20
25/25 [=====] - 1s 60ms/step - loss: 1.4855 - accuracy:
0.4643 - val_loss: 5.5863 - val_accuracy: 0.2100
Epoch 9/20
25/25 [=====] - 2s 62ms/step - loss: 1.3332 - accuracy:
0.5031 - val_loss: 4.2907 - val_accuracy: 0.2050
Epoch 10/20
25/25 [=====] - 1s 60ms/step - loss: 1.3956 - accuracy:
0.4906 - val_loss: 11.5004 - val_accuracy: 0.1200
Epoch 11/20
25/25 [=====] - 2s 60ms/step - loss: 1.2097 - accuracy:
0.5607 - val_loss: 4.7443 - val_accuracy: 0.2000
Epoch 12/20
25/25 [=====] - 2s 61ms/step - loss: 1.2549 - accuracy:
0.5544 - val_loss: 6.8092 - val_accuracy: 0.2300
Epoch 13/20
25/25 [=====] - 2s 61ms/step - loss: 1.1238 - accuracy:
0.5582 - val_loss: 17.0100 - val_accuracy: 0.1650
Epoch 14/20
25/25 [=====] - 2s 60ms/step - loss: 1.0561 - accuracy:
0.6183 - val_loss: 12.6391 - val_accuracy: 0.1950
Epoch 15/20
25/25 [=====] - 2s 61ms/step - loss: 1.0557 - accuracy:
0.6095 - val_loss: 20.1991 - val_accuracy: 0.1300
Epoch 16/20
25/25 [=====] - 2s 61ms/step - loss: 1.1104 - accuracy:
0.6120 - val_loss: 20.2748 - val_accuracy: 0.0800
Epoch 17/20
25/25 [=====] - 2s 61ms/step - loss: 0.9231 - accuracy:
0.6658 - val_loss: 7.3416 - val_accuracy: 0.1950
Epoch 18/20
25/25 [=====] - 2s 65ms/step - loss: 0.7418 - accuracy:
0.7146 - val_loss: 10.3693 - val_accuracy: 0.2300
Epoch 19/20
25/25 [=====] - 2s 64ms/step - loss: 0.8063 - accuracy:
0.7121 - val_loss: 8.4882 - val_accuracy: 0.2800
Epoch 20/20
25/25 [=====] - 2s 61ms/step - loss: 0.7494 - accuracy:
0.7159 - val_loss: 21.3820 - val_accuracy: 0.1150

```

```

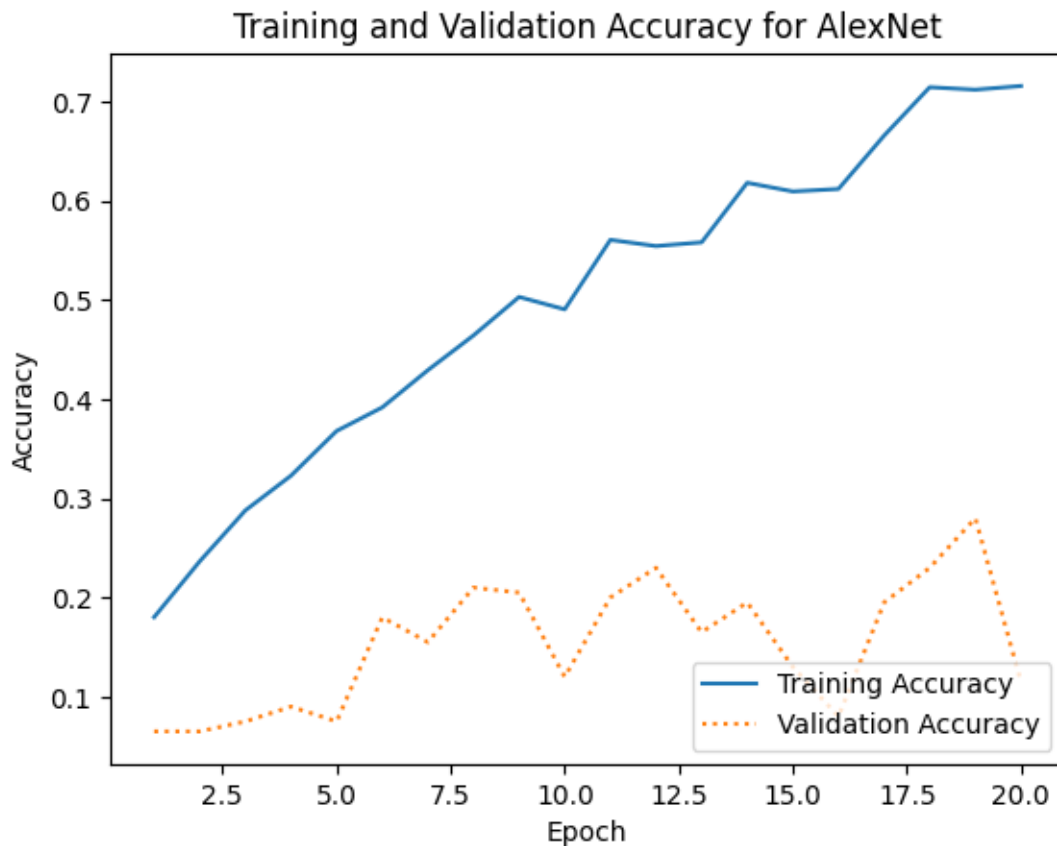
[ ]: acc = history3.history['accuracy']
      val_acc = history3.history['val_accuracy']
      epochs = range(1, len(acc) + 1)

      plt.plot(epochs, acc, '-', label='Training Accuracy')

```

```
plt.plot(epochs, val_acc, ':', label='Validation Accuracy')
plt.title('Training and Validation Accuracy for AlexNet')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
```

```
[ ]: <matplotlib.legend.Legend at 0x2c8c443d0>
```



From the training, we can see that the accuracy for different models: - Logistic regression: 0.47 - SVC: 0.41 - MLP: doesn't run - CNN model1: 0.47 - CNN model2: 0.505 - ALexNet: 0.28

We can conclude that CNN has the best performance for classifying through image data.

## 2 Conclusion

Comparing the accuracy between classification through feature data and classification through image data, we can see that feature data has better performance result. It might be because feature data might be able to capture more meaningful and relevant information about the data compared to the raw image data. It might also be less noisy, or having lower dimensional compared to raw image data, which can make it easier for the model to learn a decision boundary between classes. From

the modeling result, we can see that image data can have very high dimensionality, which can lead to overfitting or difficulty in training.

Reference: - [Work w/ Audio Data: Visualise, Classify, Recommend](#) - [Audio classification using convolutional neural networks](#) - [github: multilayer\\_perceptron](#)